

HPS Gaussian-Process Resonance Search

Combined Global Search and Observed Limits

Analysis note v4.9.16: signal extractions and staged exposure checks

Emrys Peets

6 September 2026

Scope of the result. Full 2015 and 2016 data, plus native 2021 10%, are combined over 19–250 MeV. The upper limits are pointwise, asymptotic 90% CL_s limits. The GP calculation estimates scan-wide probabilities under one archived joint stress background. It does not establish a final discovery significance or calibrate confidence-interval coverage.

1 What the combined calculation adds

At each mass, the datasets that cover that mass share one possible particle coupling. Their information is combined in a single likelihood. Membership follows the established ranges; the analysis does not choose whichever dataset gives the strongest-looking result.

The upper curve limits the allowed particle contribution at each mass. The probability curves below it ask how often the chosen background would produce something this unusual somewhere in the entire search. The look-elsewhere calculation changes those probabilities and does not automatically rescale the upper limits. The principal profiled scan has no exceedance in 1,000 direct joint experiments. Its conditional probability is unresolved; the one-sided 95% upper bound is 0.003. The Gaussian field also has no sampled exceedance. An extreme stress-centered score is not an established particle significance.

Statistic	Mass [MeV]	Local p	GP global p	Direct count
Profiled	76	6.98×10^{-20}	$<1.50 \times 10^{-5}$	0/1000
Fixed	76	3.67×10^{-13}	$<1.50 \times 10^{-5}$	0/1000

Table 1: Principal minimum-local-p ordering under the common stress background. Less-than signs are one-sided 95% Monte Carlo upper bounds, not measured zero probabilities. Fixed-background results are a separate diagnostic method.

Mass [MeV]	Active data
19–38	2015 full
39–49	2015 full + 2016 full
50–90	2015 full + 2016 full + 2021 10%
91–180	2016 full + 2021 10%
181–250	2021 10%

Table 2: The complete search uses 1 MeV spacing and established dataset ranges.

2 Full combined search

Combined HPS search over the full mass region

2015 full + 2016 full + 2021 10% | 19-250 MeV | 1 MeV grid

Global probabilities condition on one archived joint stress background.




Figure 1: The complete 232-point search: the profiled observed upper limit above local and global probabilities. The activity strip fixes which data enter each mass. The upper curve includes the visible branching correction once above the dimuon threshold. Lower panels condition on one common joint stress background. Gray local probabilities use the dense observed root and the asymptotic convention; the Gaussian-response local rule assigns one to nonpositive raw roots. Local triangles mark values below the 10^{-8} display floor. Global triangles are one-sided 95% Monte Carlo upper bounds for zero simulated tails, separated from probability estimates. Green points check the principal ordering; their error bars are central 95% intervals. The two global orderings are different tests. Expected-limit bands are not shown.

3 Illustrative scan of deficits

A negative auxiliary signal amplitude describes missing events with the signal-template shape; it is not a physical negative rate or coupling. This figure mirrors the excess scan without changing its fits or limits.




Figure 2: Profiled deficit scan. Top: observed signed roots and the stress-background offset; shading marks negative fitted contributions. Middle: raw-root Gaussian reference and conditional Gaussian-response tails. Both assign one to nonnegative raw roots. Local triangles mark the 10^{-8} display floor. Bottom: two separate global orderings; triangles mark one-sided 95% zero-count bounds. Direct error bars are central 95% intervals.

For $z = (r - a)/s$, the conditional local rule is $p^- = \Phi(z)$ when $r < 0$, and one otherwise. The principal scan score is $T^- = \max_{m:r_m < 0}(-z_m)$, with $-\infty$ for an empty set. The separate raw-depth statistic is $D^- = \max_m \max(0, -r_m)$.

The deepest raw profiled deficit is at 72 MeV, with $r = -3.490$ and an uncalibrated local reference of 2.41×10^{-4} . The largest stress-centered deficit instead occurs at 83 MeV: $r = -0.676$, $a = +7.707$ and $s = 0.983$. Its raw reference is 0.2495, whereas the conditional local value is 7.18×10^{-18} . This contrast again comes mainly from the stress offset.

The latter threshold has 0/200,000 GP and 0/1,000 direct exceedances. The one-sided 95% global upper bounds are 1.50×10^{-5} and 0.003, respectively. Every simulated raw-depth maximum exceeds the observed deficit. These are the same realizations used for the excess scan; no new fits or independent toys were needed.

These illustrative tails follow the excess scan and do not adjust for choosing directions, methods or orderings. The unresolved tails establish neither a particle signal nor physical background validity.

4 Selected signal extractions and staged exposure checks

What the displays establish. The largest observed combined excess is at 66 MeV. All three years prefer a positive fitted signal there, although each estimate has substantial uncertainty. The next peak in the complete search is 21 MeV, where only 2015 contributes. The next leading peak covered by several datasets is 92 MeV, where 2016 and 2021 prefer noticeably different rates. These are useful locations for follow-up; the selected plots alone do not establish a particle or a failure of the background model.

The strongest raw combined deficit is at 72 MeV. Showing it alongside the excesses matters: a background interpolation error can produce both positive and negative residuals, while a positive narrow signal does not by itself explain a nearby missing-event feature. Statistical fluctuations can also produce such patterns. An independent data increment and predictive background checks can distinguish these possibilities more directly than enlarging a selected peak display.

Which locations are shown

Selection uses the dense *observed profiled* signed root, whose ordering matches its one-sided asymptotic local reference on the positive branch. Select positive local maxima, ordered by decreasing root, with a separation greater than 2.25 times the larger resolution at the two coordinates. For the combined rule use the largest active resolution. Endpoints remain eligible and are identified. The stress-centered ordering is discussed separately; it does not choose the leading extraction panels.

Data	First excess	Second excess	Deepest deficit
2015 full	51 (+3.14)	21 (+2.52)	19 (-3.21)
2016 full	90 (+3.42)	117 (+3.28)	102 (-4.50)
2021 10%	78 (+2.81)	65 (+2.40)	71 (-4.02)
Full combined union	66 (+2.76)	21 (+2.52)	72 (-3.49)

Table 3: Selected mass in MeV, followed by the signed root in parentheses. These are selected local fit summaries, not global significances. The 19 MeV deficit is the 2015 search endpoint. The additional combined 92 MeV display has $r = +2.42$ and uses 2016 plus 2021.

How to read the figures

The top row shows event counts, the GP mean before profiling, the background fitted together with the signal, and their sum. The lower row subtracts that fitted background and overlays the fitted signal. Black bars show counting uncertainty only. The gray envelope is a zero-centered width guide from the original correlated GP constraint, projected into the display bins. It does not depict that constraint’s center, an error band on the fitted background, or the full residual covariance. These residuals are not independent significance measurements.

Display bins are whole original bins grouped to approximately half the mass resolution, with a fixed histogram-edge origin. The likelihood retains every native fit bin. Multi-year sums use a common 1.25 MeV grid restricted to the overlap of the fitted windows. Each year keeps its own resolution, acceptance and count-to-coupling conversion in the likelihood. No fitted curve is extended outside its fit window.

The leading combined excess: 66 MeV





Figure 3: The 66 MeV common-amplitude fit, shown separately for each year and as a display sum. The shared signed amplitude is $\hat{\epsilon}^2 = 4.55 \times 10^{-6}$ and $r = 2.76$. Every panel uses the same amplitude, with the appropriate year-specific yield and resolution. Summed panels omit boundary pieces; the likelihood root and quoted fitted yields use the full native windows.

A second location with multiple datasets: 92 MeV

Second leading peak with multiple datasets: 92 MeV

Common-amplitude profile fit | signed root $r = +2.42$ | fitted amplitude = $+2.97 \times 10^{-6}$



All panels use the same fitted amplitude. The summed panel uses only common whole bins; the likelihood uses the complete native windows.

Bars: counting error only. Gray shading: zero-centered GP constraint width; not fitted-background error or total residual uncertainty.

Figure 4: The 92 MeV fit uses 2016 and 2021. The common amplitude is $\hat{\epsilon}^2 = 2.97 \times 10^{-6}$ and $r = 2.42$. Both separate fits are positive, but their preferred rates differ (Fig. 10). A persistence illustration at this common rate is a testable assumption.

The deepest combined observed deficit: 72 MeV

Deepest combined observed deficit: 72 MeV

Common-amplitude profile fit | signed root $r = -3.49$ | fitted amplitude $= -4.88 \times 10^{-6}$

● Observed data - - - Background fitted with signal ■ GP constraint SD (width guide)
 GP mean before profiling — S+B (top); fitted S (bottom)





Figure 5: The deepest raw combined deficit has $r = -3.49$ and a signed auxiliary amplitude $\tilde{\epsilon}^2 = -4.88 \times 10^{-6}$. All three separate amplitudes are negative, with 2021 providing the strongest individual contribution. Negative templates describe missing events relative to the fitted background, not a physical negative coupling.

Individual signal extractions: 2015 full

The 51 MeV point is the largest observed profiled excess in 2015. The 21 MeV point is also the second leading peak of the full combined union, because 2015 is the only active dataset at this mass. It is only 2 MeV above the search endpoint; the 19 MeV deficit in Fig. 9 makes boundary and background-shape controls particularly relevant. More 2021 data cannot test this 21 MeV location.

2015 full: two leading observed excesses

Separate single-dataset fits | observed-profiled ranking | display bins $\approx$ half the mass resolution



Selected after scanning. Local signed roots describe the fitted templates; these panels do not establish a global particle significance.

Bars: counting error only. Gray shading: zero-centered GP constraint width; not fitted-background error or total residual uncertainty.

Figure 6: The two leading separated observed profiled excesses in 2015 full. Each panel is an independent single-dataset fit. Binning is chosen from the resolution and original histogram origin, not from the observed residual shape.

The lower panels are background-subtracted displays conditioned on the fitted template. A visible peak here should be assessed together with the background constraint, the fitted amplitude uncertainty, the scan selection and independent-data checks.

Individual signal extractions: 2016 full

The two leading individual 2016 peaks are at 90 and 117 MeV. Their positions differ from the leading combined mass because a common coupling also has to describe the other active dataset contributions. These panels retain the inherited 2016 source-fit waiver, development overlap and numerical qualifications; cleaner presentation does not resolve them.




Figure 7: The two leading separated observed profiled excesses in 2016 full. Each panel is an independent single-dataset fit. Binning is chosen from the resolution and original histogram origin, not from the observed residual shape.


The lower panels are background-subtracted displays conditioned on the fitted template. A visible peak here should be assessed together with the background constraint, the fitted amplitude uncertainty, the scan selection and independent-data checks.

Individual signal extractions: 2021 10%

The released 2021 sample has leading individual peaks at 78 and 65 MeV. These are near, but not interchangeable with, the combined coordinates. Changing the mass or window after seeing a new sample would create another selection. Freeze these individual locations as a declared secondary family if they are followed at the 30% checkpoint.

2021 10%: two leading observed excesses

Separate single-dataset fits | observed-profiled ranking | display bins $\approx$ half the mass resolution



Selected after scanning. Local signed roots describe the fitted templates; these panels do not establish a global particle significance.

Bars: counting error only. Gray shading: zero-centered GP constraint width; not fitted-background error or total residual uncertainty.


Figure 8: The two leading separated observed profiled excesses in 2021 10%. Each panel is an independent single-dataset fit. Binning is chosen from the resolution and original histogram origin, not from the observed residual shape.

The lower panels are background-subtracted displays conditioned on the fitted template. A visible peak here should be assessed together with the background constraint, the fitted amplitude uncertainty, the scan selection and independent-data checks.

Individual deficits as a background diagnostic

Deepest observed deficit in each dataset

Independent fitted amplitudes | same profile construction and resolution-based binning



The negative signal template is an auxiliary diagnostic of missing events; it is not a physical negative event rate.


Bars: counting error only. Gray shading: zero-centered GP constraint width; not fitted-background error or total residual uncertainty.

Figure 9: Deepest observed negative roots in 2015, 2016 and 2021. The 2015 point lies at the scan endpoint. The 2016 and 2021 minima are at 102 and 71 MeV, respectively. The negative branch provides a diagnostic with the same narrow template; no physical negative signal rate or calibrated deficit discovery is inferred.

Do separate years support the same rate?

Do the years prefer a compatible signal rate?

Points: separate profiled fits | Red line and shading: shared-amplitude fit $\pm$ local curvature SD



At 66 MeV all three fitted amplitudes are positive. At 92 MeV the preferred rates differ; at 76 MeV the years pull in opposite directions.

ΔD compares independent amplitudes with one shared amplitude at a selected mass. No post-selection or global compatibility probability is claimed.

Figure 10: Independent signed amplitude estimates and local curvature standard errors, compared with the common fit. Define $\Delta D = 2[\text{NLL}_{\text{common}} - \sum_d \text{NLL}_{d,\text{free}}]$. The number of additional amplitude parameters is one fewer than the number of datasets. These masses were selected from the data, so the displayed likelihood losses are conditional descriptive comparisons; no post-selection compatibility probabilities are quoted.


Why the most extreme stress-centered tail need not look like a peak

The stress-background centering changes the reference distribution of the statistic; it does not change the observed fitted signal. At 76 MeV the combined raw root is only $+0.166$, but the generating stress construction has Asimov root $a = -8.700$ and response width $s = 0.979$. Subtracting this large negative offset produces a standardized value of 9.05 . At 83 MeV the observed raw root is -0.676 , with $a = +7.707$ and $s = 0.983$, giving a standardized deficit depth of 8.53 .

The very small conditional tails therefore ask about departure from those particular stress constructions. They do not mean that the data contain a correspondingly large positive or negative fitted signal. At 76 MeV, 2016 prefers a negative amplitude while 2021 prefers a positive one. The common-versus-independent likelihood loss is $\Delta D = 9.40$ for two additional amplitudes. This provides another reason to avoid presenting the stress-centered tail as a coherent resonance.

Extreme stress-centered tails can have small fitted signals

76 MeV: stress offset $a = -8.70$ | 83 MeV: stress offset $a = +7.71$



Centering changes the probability assigned to a fit. It does not turn the small red template into a larger fitted signal.

Bars: counting error only. Gray shading: zero-centered GP constraint width; not fitted-background error or total residual uncertainty.

Figure 11: Common-window sums at the two principal stress-centered extrema. The small fitted signal templates are consistent with their small raw roots. Large stress offsets cause the extreme centered-tail interpretation. These sums have no independent fit.

The asymptotic profiled method is not invalid merely because calibration changes a probability or an observed limit. Its use in a final analysis requires qualification of the background family, the prediction uncertainty, signal response and the relevant sampling distributions. The present displays do not measure expected sensitivity. A weaker calibrated observed limit can reflect a correction to background behavior rather than information being removed from the data.

A 30% checkpoint before the full 2021 sample

Use the released 10% to define the hypotheses. Treat an additional, disjoint 20% as the next independent test of those fixed locations and rates. Only then inspect the cumulative 30% combination. This is more informative than a cumulative display alone, because the latter still contains the fluctuation that selected the peak. The same original-event membership and processing must be documented for every increment.

Let N_{10} be the observed first sample, B_{10} its background fitted in the selected common-amplitude model, and S_{10} the signal template at that fitted rate. Figure 12 compares two illustrative conditional means for a total exposure f times the original 10%:

$$\begin{aligned} \text{background only in added data: } & N_{10} + (f - 1)B_{10}, \\ \text{selected rate persists: } & N_{10} + (f - 1)(B_{10} + S_{10}). \end{aligned}$$

The independent additional 20% has means $2B_{10}$ or $2(B_{10} + S_{10})$. For the conditional cumulative 30% and 100% views the background-only counting variances from the future sample are $2B_{10}$ and $9B_{10}$, respectively. The observed first 10% is held fixed. These figures are mean scenarios, not new observations, toy experiments or discovery projections. They omit uncertainty in the assumed background and selected signal rate. The 92 MeV common-rate assumption is less consistent with the individual estimates than the 66 MeV assumption.

Mass [MeV]	10%	new 20%	30% total	100% total
66	14.6	29.2	43.7	145.8
92	13.6	27.2	40.9	136.2

Table 4: Illustrative mean signal yields in thousands of events in the complete 2021 fitted window if the selected common-fit rate persists. These are template yields, not the raw total counts or measured future yields. Display sums may omit boundary pieces.

A statistics-dominated precision reference can be computed without toys. With the original per-year signal yield vector u_d , GP constraint covariance C_d and GP mean b_d , define

$$I_d = u_d^T [\text{diag}(b_d) + C_d]^{-1} u_d, \quad I(f) = I_{2015} + I_{2016} + f I_{2021}.$$

An absent dataset has $I_d = 0$. This local background-Asimov reference follows the usual information-scaling argument [2] and assumes that the *entire* 2021 count covariance scales as f , with unchanged acceptance and resolution. Fractional systematic effects or background bias need not improve this way. The factor $\sqrt{I(f)/I(1)}$ below is a precision illustration, not a multiplier for the observed combined root or a forecast global significance.


Mass [MeV]	Current 2021 information share	30% gain	100% gain
66		55.3%	1.45
92		82.4%	1.63

Table 5: Combined precision gain with 2015 and 2016 held fixed, under the stated covariance-scaling assumption. For 2021 alone the corresponding factors would be $\sqrt{3}$ and $\sqrt{10}$. There is no 2021 gain at 21 MeV.

Conditional exposure displays: 10%, new 20%, 30%, 100%

A 30% checkpoint before opening the full 2021 sample

Illustrative persistence of the selected common-fit rate at 66 and 92 MeV | 2015 and 2016 stay fixed



Future panels are conditional expectations. Cumulative panels retain the actual first 10%; y-axis ranges differ to show the change in event yield.

The 92 MeV rates differ between years. These selected-rate examples omit background-model uncertainty; future data and global significances are not supplied.

Figure 12: Background-subtracted 2021 views at fixed 66 and 92 MeV hypotheses. Only the first column contains observed points. Future curves keep the first sample fixed where included and add the mean of either hypothesis. Gray envelopes show only the background counting standard deviation of the added sample. They exclude fitted-background uncertainty and do not imply independent residual bins after a future refit. The vertical scales differ.

What would make the next comparison decisive?

First, define a disjoint additional 20% with the same selection, calibrations, mass resolution and normalization convention. Freeze the 66 and 92 MeV primary locations, and state in advance whether the individual 2021 locations at 65 and 78 MeV and the 71/72 MeV deficit checks are a secondary family. Do not reselect the mass from the new spectrum when reporting the fixed-location validation.

Fit the additional 20% on its own before combining it with the old 10% or the 2015/2016 samples. Compare its fitted rate, width behavior and background residuals with the stored predictions. A stable signal should reproduce a compatible rate, not merely increase the height of a cumulative display. Rate agreement alone would still leave detector effects or common background structure to investigate.

Use predeclared sideband and predictive checks to test the background interpolation and relevant detector/selection effects, including both positive and negative residuals. The 2016 background-source transition and existing numerical/source qualifications remain part of that work. Model choices must not be selected by whichever one makes the observed peak more impressive.

The 30% and 100% cumulative looks share events. Before treating either as a formal discovery or exclusion result, declare the family of mass, direction and exposure tests and the treatment of sequential looks. Rebuild the appropriate signal and background sampling distributions for the new exposure and background family. The existing GP global-tail bank describes the current released sample and cannot calibrate a different exposure.

For an efficient follow-on calculation, begin with ten complete pilot spectra per declared generating hypothesis, retaining each same full spectrum throughout its mass scan. Use new deterministic seed namespaces and nonoverlapping toy IDs. Time the pilot, test exact-fit and approximation closure, then scale in restartable batches only if the validated cost is acceptable. Background-only scans address global discovery tails; calibrated limits additionally require signal-plus-background hypotheses. No new toys were necessary for the present extraction and mean-scenario displays.

Display and reconstruction checks

Fifteen selected likelihoods were reconstructed with the archived dense solver and exact prediction hashes. There are 24 dataset components. Every background-plus-signal vector closes to the original fitted expectation; all expected Poisson counts stay positive. Every grouping matrix contains only zero or one, with no fractional reassignment or reused bin within a panel. The quoted roots and upper endpoints reproduce the frozen dense scan; display grouping never changes a fit.

Mass [MeV]	Dataset	Counts retained in sum	Signal template retained
66	2015	60.1%	87.7%
66	2016	83.6%	96.1%
66	2021	100.0%	100.0%
92	2016	57.8%	83.3%
92	2021	93.8%	98.8%
72	2015	48.0%	77.3%
72	2016	68.0%	88.4%
72	2021	93.1%	98.3%

Table 6: Common-window summed-display retention relative to the complete native fitted window of each year. The full likelihood, amplitude estimates and total-window yield tables use all fit bins. Individual panels have their own resolution-based grouping.

The reusable PDF and PNG figures, native and rebinned arrays, exact grouping maps, exposure tables, build scripts and independent HEP audit are retained in

`study_results/v4p9p16_presentation_extractions_20260906/`

The prior studies were committed, pushed and merged in PR 67; all 4,778 published file blobs were verified on the merged branch. This presentation extension has a separate artifact manifest. The original study inputs and outputs remain frozen.

5 One shared-coupling likelihood

Let $\mathcal{D}(m)$ be the predeclared active dataset set. The model combines Poisson bins and independent Gaussian background constraints:

$$\mathcal{L}_m(\epsilon^2, \theta) = \prod_{d \in \mathcal{D}(m)} \prod_{i \in W_d(m)} \text{Pois}(n_{di} \mid b_{di} + (L_d \theta_d)_i + \epsilon^2 S_{di}(m)) \prod_{d \in \mathcal{D}(m)} e^{-\theta_d^\top \theta_d / 2}.$$

$S_{di} = \text{BR}_{ee} T_{di}$ is the visible yield per unit physical coupling, where T_{di} is the stored unit-electron-branching template. The fitted coordinate is $\epsilon_{\text{eff}}^2 = \epsilon_{\text{phys}}^2 \text{BR}_{ee}$; the display restores $\epsilon_{\text{phys}}^2 = \epsilon_{\text{eff}}^2 / \text{BR}_{ee}$. The implementation rescales this coordinate into a total window amplitude. No independent signal strength is fitted in each year.

An auxiliary unrestricted amplitude supplies the signed likelihood-ratio root,

$$r_m = \text{sign}(\hat{A}_m) \sqrt{2 \left[\text{NLL}_m(A=0) - \text{NLL}_m(\hat{A}_m) \right]}.$$

Here $\text{NLL}_m(A)$ is the negative log likelihood minimized over the background nuisances at fixed A . Poisson expectations remain positive. Physical discovery and upper-limit tests retain the nonnegative signal boundary; the auxiliary negative branch is needed to model the signed field.

Complete joint experiments

Ten pilot and 1,000 separate validation spectra per year are reused from v4.9.14/v4.9.15. The year-specific random streams are distinct. Pairing equal row indices gives 1,000 joint experiments, not 3,000. Each experiment supplies exactly the same full spectrum at every mass in which its dataset participates.

Kernels, templates, support and exclusion masks remain frozen. Every toy retrains its log-count background prediction and count-dependent training errors. All multi-dataset intervals receive new common-amplitude fits. Only the single-dataset intervals reuse existing likelihood-root columns.

The generating means are the archived smooth stress spectra. Event-level independence of the 2016 development input from its full observed sample has not been established. Independence here describes the conditional toy model; omitted experimental systematic correlations have not been newly quantified.

One response basis across boundaries

Following Ananiev and Read [1], perturb each data bin by its Poisson standard deviation. This extension retains the nonzero offset and response width:

$$a_m = r_m(B), \quad D_{im} = r_m(B + \sqrt{B_i} e_i) - a_m, \quad C = D^\top D, \quad s_m = \sqrt{C_{mm}}, \quad K_{mn} = \frac{C_{mn}}{s_m s_n}.$$

The basis has 484, 720 and 422 bins from 2015, 2016 and 2021. An inactive dataset has exactly zero response at that mass. This retains correlations across boundaries that share data; separately sampled segment maxima are not joined.

For each statistic, 200,000 vectors $Z \sim N(0, K)$ generate $r_m^* = a_m + s_m Z_m$. The principal maximum is $T = \max_{m: r_m^* > 0} Z_m$ (or $-\infty$ for an empty set); the secondary maximum is $R = \max_m \max(0, r_m^*)$. Their respective tails are compared with the independent joint Poisson scans.

6 Offsets and widths of the combined statistic




Figure 13: Asimov-response predictions compared with 1,000 joint Poisson scans. Top panels show the offset in the signed root; bottom panels show its spread. Vertical lines mark membership changes. Agreement tests the approximation under the generating means, without establishing their physical validity.

Statistic	Standardized mean range	Spread range	Flags
Profiled	-0.079 to +0.076	0.957 to 1.074	0
Fixed	-0.093 to +0.080	0.943 to 1.050	0

Table 7: Normality flags use a Holm adjustment across 232 masses within each method. Non-rejection has finite power and cannot certify extreme Gaussian tails.

Centering, scaling and covariance come from deterministic bin perturbations. Validation scans do not tune them. Profiled and fixed-background statistics use identical spectra but can have different sampling widths.

Large oscillating offsets can make an ordinary observed value unusual relative to a particular stress construction. A small probability can diagnose that construction without identifying a particle. The inherited 2016 numerical exception, source-fit waiver, development overlap and 75–85 MeV source transition remain unresolved by this extension.

7 Correlations and the look-elsewhere check




Figure 14: Response correlations across the union grid. Dotted lines show membership changes. Shared datasets retain correlations across boundaries; regions with no common dataset have zero response covariance in this model.




Figure 15: Principal scan-maximum tails. Direct-toy bands are pointwise central 95% binomial intervals, not expected-limit bands or simultaneous coverage bands. Triangles mark the upper endpoint of the central interval at the first unresolved direct tail. More GP draws reduce sampling error within the approximation, without replacing direct validation.

Profiled: the maximum-distribution KS diagnostic gives nominal $p=0.8625$ for the principal ordering and 0.6389 for the raw ordering. The correlation RMS difference is 0.031.

Fixed: the maximum-distribution KS diagnostic gives nominal $p=0.9974$ for the principal ordering and 0.2034 for the raw ordering. The correlation RMS difference is 0.030.

These diagnostics compare simulation methods, not the data significance. Two methods are checked per ordering. Direct Poisson counts calibrate the declared score under the generating spectrum without requiring a jointly Gaussian field; the GP estimate has that additional approximation.

8 Representative probabilities

For positive raw roots, the principal local rule is $p_m = \bar{\Phi}[(r_m - a_m)/s_m]$; otherwise $p_m = 1$. Its global curve counts how often any of the 232 masses reaches that local threshold. The separate raw ordering scans $\max_m \max(0, r_m)$. These tests cannot be selected according to which gives the smaller probability.

Mass	Active years	Asymp. local	GP local	GP global	Direct	Raw global
30	2015	0.1291	0.01	0.5479	543/1000	1
65	2015+2016+2021	0.0052	1	1	1000/1000	1
66	2015+2016+2021	0.0029	1	1	1000/1000	1
76	2015+2016+2021	0.4342	6.98×10^{-20}	$<1.50 \times 10^{-5}$	0/1000	1
120	2016+2021	0.0795	0.0598	0.9813	975/1000	1
220	2021	0.5	1	1	1000/1000	1

Table 8: Profiled representative masses: the principal and raw extrema, plus 30, 65, 120 and 220 MeV. GP local uses the common-background Gaussian response. Both global columns use the full union at the threshold observed at that mass. Raw global is the separate raw-peak test.

Statistic	Mass	Raw r	Offset a	Width s	$(r - a)/s$
Profiled	76	+0.166	-8.700	0.979	+9.053
Fixed	76	+0.175	-15.348	2.164	+7.173

Table 9: Principal peak decomposition. The final column is a conditional score, not an independently established discovery significance.

At the profiled principal peak, the unadjusted asymptotic local probability is 0.4342, while centering and scaling against the stress background gives 6.98×10^{-20} . The observed raw root is close to zero; the extreme centered score arises mainly because the stress background predicts a large negative root.

At 74 and 75 MeV, the centered scores are 10.20 and 10.38, but the raw roots are negative. The declared boundary therefore assigns local $p=1$ there. The next positive raw root becomes the principal peak. This abrupt change comes from the signal boundary together with the stress offset; it is not a numerical sign flip.

The raw maximum occurs at 66 MeV. All 200,000 GP fields and all 1,000 direct experiments exceed it. These differences reflect the specified orderings and background offsets; they are not a free choice of significance. The raw maximum tests positive peaks; it does not assess full-spectrum goodness of fit.

Statistic	Full 1 MeV grid	First 2 MeV subgrid	Second subgrid
Profiled	$<1.50 \times 10^{-5}$	$<1.50 \times 10^{-5}$	$<1.50 \times 10^{-5}$
Fixed	$<1.50 \times 10^{-5}$	$<1.50 \times 10^{-5}$	$<1.50 \times 10^{-5}$

Table 10: GP probabilities at the same original peak threshold. Zero-count bounds do not establish fine-tail convergence.

Subgrids remove alternate masses at the original observed threshold. They test finite-grid dependence, not convergence to a continuous search. Probabilities include neither a choice between methods nor uncounted searches over widths, kernels or other analysis choices.

9 Observed limits and the v4.9.12 comparison




Figure 16: The pointwise upper-limit construction with the current dense and earlier cached solvers. The lower panel shows their relative difference. Both curves have the same membership and branching correction. Their difference is not a look-elsewhere adjustment or a sensitivity measurement.

The largest absolute relative change from the corresponding v4.9.12 limit is 2.502%. The largest difference in the nonnegative discovery root is 0.0267. Fresh dense fits reproduce v4.9.13 dense references where the same scope exists. The old optimizer does not select a preferred answer. Above 211.317 MeV, displayed limits include $1 + \sqrt{1 - 4m_\mu^2/m^2}(1 + 2m_\mu^2/m^2)$ once; the CSV retains the electron-channel values.

The upper limit inverts the bounded, piecewise-asymptotic CL_s statistic at each mass [2]. The background-only joint scans test the discovery-score approximation. They do not supply the signal-plus-background distributions needed to calibrate the full CL_s curve. A GP global probability is not a replacement CL_s tail.

The v4.9.13 study contains conditional calibrated individual limits and the common all-three interval. It has no calibrated pairwise endpoints over the remaining union. Splicing different calibrations would conceal a change of construction. A fully calibrated union limit requires additional declared signal-plus-background work; this note supplies the complete observed asymptotic curve.

10 Numerical checks, provenance and continuation

All 2736 numerical and product-identity checks passed. There are 142 newly fitted joint coordinates and 90 reused single-dataset coordinates. Summed coordinate execution took 24.1 minutes with one worker and one BLAS thread. The exact backend was retained at 1 new coordinates (41 MeV). The largest paired pilot-root difference is 2.10×10^{-5} , the largest complete-response relative vector error is 4.73×10^{-6} , and the largest sentinel-correlation error is 4.65×10^{-7} . No paired pilot changed its raw-positive classification.

Every new joint mass has ten paired exact pilot fits, an exact baseline and a fixed bin-response stencil. Complete exact responses are retained at 39, 49, 50, 90, 91 and 180 MeV. The accelerator must pass root, centered-response, width and correlation gates. Otherwise all ensembles at that coordinate use the exact backend. Failed exact references halt; no toy or hypothesis is removed.

The v4.9.14 and v4.9.15 inputs are checked against their full frozen manifests. The v4.9.12 comparison table is checked against its specific frozen-manifest entry. The isolated derivative is

```
study_results/v4p9p16_combined_global_20260906/
```

It retains spectra, joint root vectors, limits, covariance, GP maxima, probability tables and per-mass audits. The independent HEP review is under `review/`. The original manifest remains unchanged. The deficit extension and this revised report have a separate manifest under

```
study_results/v4p9p16_deficit_extension_20260906/
```

The independent deficit audit passed 939 conditions.

From the repository root, verify the completed numerical products with

```
python3 -B \  
study_results/v4p9p16_combined_global_20260906/verify_combined.py
```

Preserve the finished version before reanalysis or new generation. `NEXT_STEPS.md` gives resume commands, new-seed requirements and the additional work for calibrated limits.

The next scientific priority is to qualify the joint background family with predeclared predictive controls. Direct tail precision, finer grids and signal injections answer separate questions. Additional GP sampling alone cannot establish physical background validity, expected sensitivity or interval coverage.

References

- [1] V. Ananiev and A. L. Read, *Gaussian Process-based calculation of look-elsewhere trials factor*, JINST **18** (2023) P05041. arXiv:2206.12328v3.
- [2] G. Cowan, K. Cranmer, E. Gross and O. Vitells, *Asymptotic formulae for likelihood-based tests of new physics*, Eur. Phys. J. C **71** (2011) 1554. arXiv:1007.1727.